

In [1]:

```
1 import pandas as pd
2 import numpy as np
3 from matplotlib import pyplot as plt
4 import seaborn as sns
5 from IPython.display import display
6 from collections import Counter
7 from imblearn.over_sampling import SMOTE
8 from sklearn.model_selection import train_test_split
9 from sklearn.metrics import accuracy_score, classification_report, roc_curve, roc_auc_score
10 from sklearn.ensemble import ExtraTreesClassifier
11 from sklearn.linear_model import LogisticRegression
12 from sklearn.ensemble import GradientBoostingClassifier
13 from sklearn.ensemble import RandomForestClassifier
14 import matplotlib.pyplot as plt
15 import scikitplot as skplt
16 from IPython.display import display
17 #from pycaret .classification import *
```

In [2]:

```
1 dtypes = { 'Regiao': 'object',
2             'UF': 'object',
3             'CNAE': 'object',
4             'Atendida': 'bool',
5             'CodAssunto': 'object',
6             'SexoConsumidor': 'object',
7             'FaixaEtaria': 'object',
8             'CEP': 'object',
9             'InscritoDAU': 'bool'}
```

In [3]:

```
1 df_ml1 = pd.read_csv(r'C:\Users\73594253368\Desktop\Curso\Datasets\Procon\dataset_trata
```

In [4]:

```
1 df_ml1 = df_ml1[['Regiao', 'UF', 'CNAE', 'Atendida', 'CodAssunto', 'SexoConsumidor', 'FaixaEtaria', 'CEP', 'InscritoDAU']]
```

In [5]:

```
1 # Este df_ml1 foi a primeira tentativa. Manteremos esse data frame para testes demonstrativos
2 # Para o ML "oficial", copiaremos esse df_ml1 para o df_ml
3 df_ml = df_ml1
```

In [6]:

1 df_ml

Out[6]:

| | Regiao | UF | CNAE | Atendida | CodAssunto | SexoConsumidor | FaixaEtaria | CE |
|-------|---------|-----|-----------|----------|------------|----------------|-------------|----------|
| 0 | Norte | RO | 6120501.0 | False | 187.0 | M | 5 | 76824042 |
| 1 | Norte | RO | 3514000.0 | False | 185.0 | M | 4 | 76824322 |
| 2 | Norte | RO | 8599604.0 | True | 236.0 | M | 3 | 78932000 |
| 3 | Norte | RO | 6120501.0 | True | 187.0 | M | 5 | 78932000 |
| 4 | Norte | RO | 6493000.0 | False | 57.0 | M | 6 | 76821331 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 10514 | Sudeste | SP | 6110801.0 | True | 187.0 | F | 4 | 9617000 |
| 10515 | Norte | RO | 6143400.0 | True | 259.0 | M | 2 | 76940000 |
| 10516 | Norte | RO | 6422100.0 | False | 63.0 | F | 6 | 76990000 |
| 10517 | Norte | RO | 3514000.0 | False | 185.0 | F | 4 | 76807400 |
| 10518 | Norte | RO | 6423900.0 | True | 53.0 | F | 3 | 76806420 |

10519 rows × 9 columns

In [7]:

```
1 # Demonstração do desbalanceamento na variável "target"
2 df_ml['Atendida'].value_counts()
```

Out[7]:

```
True      6306
False     4213
Name: Atendida, dtype: int64
```

Aplicando SMOTE

Data Preparation

As variáveis preditoras mais importantes do nosso dataset são as categóricas. Assim, como etapa preparatória do SMOTE, temos que criar variáveis dummies. Testamos, antes, com SMOTE e sem dummies e, também, o tradicional dummies sem SMOTE. Igualmente testamos LabelEncoder + dummies + SMOTE. Os melhores resultados de acurácia e recall foram com o procedimento a seguir.

In [8]:

```
1 df_ml.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10519 entries, 0 to 10518
Data columns (total 9 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   Regiao                10519 non-null  object
 1   UF                    10519 non-null  object
 2   CNAE                  10519 non-null  object
 3   Atendida              10519 non-null  bool
 4   CodAssunto            10519 non-null  object
 5   SexoConsumidor        10519 non-null  object
 6   FaixaEtaria           10519 non-null  object
 7   CEP                   10519 non-null  object
 8   InscritoDAU           10519 non-null  bool
dtypes: bool(2), object(7)
memory usage: 595.9+ KB
```

In [9]:

```
1 df_ml = pd.get_dummies(df_ml[['Regiao',
2                               'UF',
3                               'CNAE',
4                               'Atendida',
5                               'CodAssunto',
6                               'SexoConsumidor',
7                               'FaixaEtaria',
8                               'CEP', 'InscritoDAU']])
```

In [10]:

```
1 df_ml.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10519 entries, 0 to 10518
Columns: 6928 entries, Atendida to CEP_9990244.0
dtypes: bool(2), uint8(6926)
memory usage: 69.5 MB
```

In [11]:

```
1 df_ml.shape
```

Out[11]:

```
(10519, 6928)
```

SMOTE após dummies

In [12]:

```

1 X = df_ml.drop(['Atendida'],axis=1)
2 y = df_ml.Atendida
3 smt = SMOTE()
4 X_os,y_os = smt.fit_sample(X,y) #os de oversampled
5 counter = Counter(y_os)
6 print(counter)

```

Counter({False: 6306, True: 6306})

In [13]:

```

1 #grafico da nova distribuição de classes
2 fig, ax = plt.subplots()
3 sns.countplot(y_os, ax=ax)
4 ax.set_title('Distribuição das Classes')
5 plt.xlabel('Classe')
6 plt.ylabel('Quantidade')
7 plt.tight_layout();
8
9 #print do balanceamento
10 print(pd.Series(y_os).value_counts())

```


D:\ANACONDA\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

warnings.warn(

True 6306

False 6306

Name: Atendida, dtype: int64



In [14]:

```

1 #Train_test_split nessa oversampled
2 #Especificamos o tamanho do test_size = 0.3 pq assim as True/False do ytreinamento e as
3 xtreinamento, xteste, ytreinamento, yteste = train_test_split(X_os, y_os, test_size = 0.3)

```


In [15]:

1 xtreinamento

Out[15]:

| | InscritoDAU | Regiao_Centro-oeste | Regiao_Nordeste | Regiao_Norte | Regiao_Sudeste | Regiao_S |
|-------|-------------|---------------------|-----------------|--------------|----------------|----------|
| 11454 | False | 1 | 0 | 0 | 0 | |
| 8765 | False | 0 | 0 | 0 | 1 | |
| 3191 | False | 0 | 0 | 0 | 1 | |
| 2792 | False | 0 | 0 | 0 | 1 | |
| 9375 | True | 0 | 0 | 0 | 1 | |
| ... | ... | ... | ... | ... | ... | |
| 12177 | True | 0 | 1 | 0 | 0 | |
| 8964 | True | 0 | 0 | 0 | 1 | |
| 4682 | False | 0 | 0 | 0 | 1 | |
| 5278 | False | 1 | 0 | 0 | 0 | |
| 7598 | False | 0 | 0 | 0 | 1 | |

8828 rows × 6927 columns




In [16]:

1 xteste

Out[16]:

| | InscritoDAU | Regiao_Centro-oeste | Regiao_Nordeste | Regiao_Norte | Regiao_Sudeste | Regiao_Su |
|------|-------------|---------------------|-----------------|--------------|----------------|-----------|
| 6594 | False | 0 | 0 | 0 | 0 | |
| 6068 | False | 0 | 0 | 0 | 1 | |
| 3782 | False | 0 | 0 | 1 | 0 | |
| 1772 | False | 0 | 0 | 0 | 0 | |
| 289 | False | 0 | 0 | 0 | 1 | |
| ... | ... | ... | ... | ... | ... | |
| 9219 | True | 0 | 0 | 1 | 0 | |
| 43 | False | 0 | 0 | 0 | 1 | |
| 4719 | False | 0 | 0 | 0 | 1 | |
| 225 | False | 0 | 0 | 0 | 1 | |
| 9295 | False | 0 | 1 | 0 | 0 | |

3784 rows × 6927 columns



In [17]:

```
1 ytreinamento.value_counts()
```

Out[17]:

```
True      4414
False     4414
Name: Atendida, dtype: int64
```


In [18]:

```
1 fig, ax = plt.subplots()
2 sns.countplot(ytreinamento, ax=ax)
3 ax.set_title('Distribuição das Classes')
4 plt.xlabel('Classe')
5 plt.ylabel('Quantidade')
6 plt.tight_layout();
7
8 #print do balanceamento
9 print(pd.Series(ytreinamento).value_counts())
```

```
True      4414
False     4414
Name: Atendida, dtype: int64
```

D:\ANACONDA\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
warnings.warn(
```



In [19]:

```
1 yteste.value_counts()
```

Out[19]:

```
True      1892
False     1892
Name: Atendida, dtype: int64
```

In [20]:

```

1 fig, ax = plt.subplots()
2 sns.countplot(yteste, ax=ax)
3 ax.set_title('Distribuição das Classes')
4 plt.xlabel('Classe')
5 plt.ylabel('Quantidade')
6 plt.tight_layout();
7
8 #print do balanceamento
9 print(pd.Series(ytreinamento).value_counts())

```


D:\ANACONDA\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

warnings.warn(

```

True      4414
False     4414
Name: Atendida, dtype: int64

```



In [21]:

```

1 #Neste ponto do notebook, as bases do "df_ml" estão balanceadas pelo SMOTE e separadas

```

PyCaret preparatório

Temos as bases balanceadas e já separadas no train_test_split. Todavia, não sabemos qual modelo de machine learning aplicar. Utilizaremos a ferramenta de automação de machine learning Pycaret apenas como guia para escolher os melhores algoritmos para implementação manual.

A documentação do PyCaret dispõe que, ao utilizar o parâmetro fix_imbalance=True, a biblioteca aplica, automaticamente, a técnica SMOTE. Dessa forma, não há necessidade de aplicar as bases balanceadas por SMOTE, às quais preparamos para o ML manual. Assim, utilizaremos, no PyCaret, a base "df_ml1". Como haverá SMOTE automático, a ml1 será semelhante à "df_ml" submetida ao SMOTE

In [22]:

```

1 #PyCaret no automático mas com fix_imbalance=True
2 #pycaret_df_ml = setup(data = ml1, target='Atendida',fix_imbalance=True)
3 #modelsml1 = compare_models()
4 #resultsm1 = pull()

```

#Resultado do PyCaret:

| | Model | Accuracy | AUC | Recall | Prec. | F1 | Kappa | MCC | TT (Sec) |
|----------|---------------------------------|----------|--------|--------|--------|--------|--------|--------|----------|
| ridge | Ridge Classifier | 0.7333 | 0.0000 | 0.7642 | 0.7873 | 0.7754 | 0.4471 | 0.4477 | 16.6980 |
| et | Extra Trees Classifier | 0.7308 | 0.7868 | 0.7998 | 0.7648 | 0.7816 | 0.4311 | 0.4325 | 33.3220 |
| rf | Random Forest Classifier | 0.7291 | 0.7867 | 0.8129 | 0.7560 | 0.7833 | 0.4231 | 0.4253 | 23.5580 |
| lr | Logistic Regression | 0.7250 | 0.7892 | 0.7380 | 0.7918 | 0.7638 | 0.4355 | 0.4374 | 31.5750 |
| svm | SVM - Linear Kernel | 0.7229 | 0.0000 | 0.7315 | 0.7974 | 0.7593 | 0.4327 | 0.4393 | 16.7670 |
| dt | Decision Tree Classifier | 0.7206 | 0.7016 | 0.7967 | 0.7538 | 0.7745 | 0.4081 | 0.4094 | 15.5360 |
| lightgbm | Light Gradient Boosting Machine | 0.7201 | 0.7718 | 0.7649 | 0.7696 | 0.7671 | 0.4163 | 0.4166 | 16.7510 |
| gbc | Gradient Boosting Classifier | 0.7027 | 0.7645 | 0.7344 | 0.7637 | 0.7486 | 0.3852 | 0.3859 | 27.2430 |
| ada | Ada Boost Classifier | 0.6925 | 0.7554 | 0.7000 | 0.7692 | 0.7327 | 0.3726 | 0.3752 | 17.9490 |
| knn | K Neighbors Classifier | 0.6345 | 0.7270 | 0.5246 | 0.7999 | 0.6333 | 0.2993 | 0.3262 | 32.0240 |
| nb | Naive Bayes | 0.5263 | 0.5962 | 0.2563 | 0.8570 | 0.3943 | 0.1625 | 0.2440 | 14.8830 |
| lda | Linear Discriminant Analysis | 0.4105 | 0.4492 | 0.4043 | 0.4638 | 0.4320 | 0.2171 | 0.2199 | 181.5890 |
| qda | Quadratic Discriminant Analysis | 0.3917 | 0.4549 | 0.1470 | 0.6861 | 0.2419 | 0.0912 | 0.1603 | 110.5300 |

ML manual a partir dos melhores modelos que prospectamos com o PyCaret: et, rf e lr. Privilegiando o recall e, também, para variar dos modelos de árvore, colocamos, também, o lightgbm. Plotamos Matriz de Confusão e grafico de ROC e AUC.

In [23]:

```

1 #Retomamos os train_test_split a partir do oversample ("_os") que já tínhamos feito

```


In [29]:

```

1 # Modelo Extra Trees Classifier
2 et = ExtraTreesClassifier(random_state=0)
3 et = et.fit(xtreinamento, ytreinamento)
4 Train_predict_et = et.predict(xteste)
5 print("Accuracy Score:", accuracy_score(yteste, Train_predict_et))
6 print(classification_report(yteste, Train_predict_et))

```

Accuracy Score: 0.7457716701902748


| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| False | 0.70 | 0.86 | 0.77 | 1892 |
| True | 0.82 | 0.63 | 0.71 | 1892 |
| accuracy | | | 0.75 | 3784 |
| macro avg | 0.76 | 0.75 | 0.74 | 3784 |
| weighted avg | 0.76 | 0.75 | 0.74 | 3784 |

In [30]:

```


1 #Matriz de Confusão
2 skplt.metrics.plot_confusion_matrix(yteste, Train_predict_et, normalize=True)
3 plt.show()

```



In [31]:

```
1 # Curva ROC e área abaixo da curva (AUC)
2 y_pred_et = et.predict_proba(xteste)
3 et_fpr,et_tpr,threshold = roc_curve(yteste,y_pred_et[:,1])
4 auc_et = auc(et_fpr,et_tpr)
5 plt.figure(figsize=(5, 5), dpi=100)
6 plt.plot(et_fpr,et_tpr, marker='.', label='Extra Trees Classifier (auc = %0.3f)' % auc_
7 plt.xlabel('False Positive Rate -->')
8 plt.ylabel('True Positive Rate -->')
9 plt.legend()
10 plt.show()
```



In [32]:

```

1 # Modelo Logistic Regression
2 lr = LogisticRegression(random_state=0)
3 lr = lr.fit(xtreinamento, ytreinamento)
4 Train_predict_lr = lr.predict(xteste)
5 print("Accuracy Score:", accuracy_score(yteste, Train_predict_lr))
6 print(classification_report(yteste, Train_predict_lr))

```

D:\ANACONDA\lib\site-packages\sklearn\linear_model_logistic.py:762: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression (https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

Accuracy Score: 0.7618921775898521


| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| False | 0.76 | 0.77 | 0.76 | 1892 |
| True | 0.76 | 0.76 | 0.76 | 1892 |
| accuracy | | | 0.76 | 3784 |
| macro avg | 0.76 | 0.76 | 0.76 | 3784 |
| weighted avg | 0.76 | 0.76 | 0.76 | 3784 |

In [33]:

```


1 skplt.metrics.plot_confusion_matrix(yteste, Train_predict_lr, normalize=True)
2 plt.show()

```



In [34]:

```
1 Y_pred_lr=lr.decision_function(xteste)
2 logistic_fpr,logistic_tpr,threshold = roc_curve(yteste,Y_pred_lr) # Y_pred_lr do decis
3 auc_logistic = auc(logistic_fpr, logistic_tpr)
4 plt.figure(figsize=(5, 5), dpi=100)
5 plt.plot(logistic_fpr, logistic_tpr, marker='.', label='Logistic (auc = %0.3f)' % auc_1
6 plt.xlabel('False Positive Rate -->')
7 plt.ylabel('True Positive Rate -->')
8 plt.legend()
9 plt.show()
```



In [35]:

```

1 # Modelo Light Gradient Boosting Machine
2 import lightgbm
3 from lightgbm import LGBMClassifier
4 lightgbm = LGBMClassifier(random_state=0)
5 lightgbm = lightgbm.fit(xtreinamento, ytreinamento)
6 Train_predict_lightgbm = lightgbm.predict(xteste)
7 print("Accuracy Score:", accuracy_score(yteste, Train_predict_lightgbm))
8 print(classification_report(yteste, Train_predict_lightgbm))

```

Accuracy Score: 0.7418076109936576


| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| False | 0.74 | 0.75 | 0.74 | 1892 |
| True | 0.75 | 0.73 | 0.74 | 1892 |
| accuracy | | | 0.74 | 3784 |
| macro avg | 0.74 | 0.74 | 0.74 | 3784 |
| weighted avg | 0.74 | 0.74 | 0.74 | 3784 |

In [36]:

```

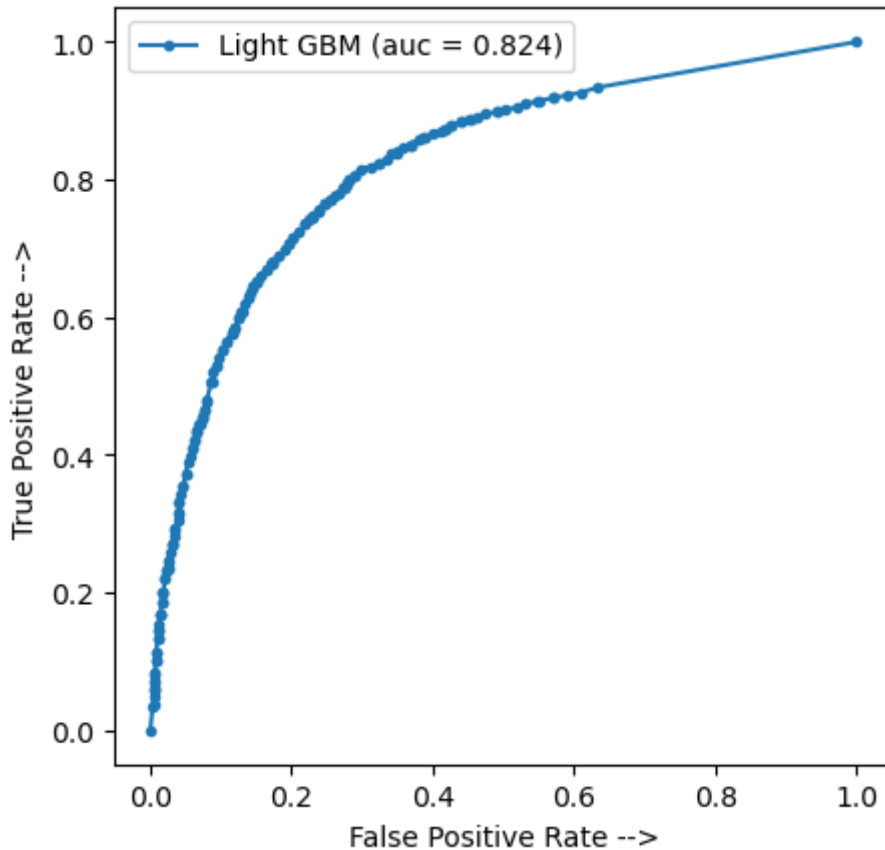
1 skplt.metrics.plot_confusion_matrix(yteste, Train_predict_lightgbm, normalize=True)
2 plt.show()

```



In [37]:

```
1 y_pred_lightgbm = lightgbm.predict_proba(xteste)
2 lightgbm_fpr,lightgbm_tpr,threshold = roc_curve(yteste,y_pred_et[:,1])
3 auc_lightgbm = auc(lightgbm_fpr,lightgbm_tpr)
4 plt.figure(figsize=(5, 5), dpi=100)
5 plt.plot(lightgbm_fpr,lightgbm_tpr, marker='.', label='Light GBM (auc = %0.3f)' % auc_1)
6 plt.xlabel('False Positive Rate -->')
7 plt.ylabel('True Positive Rate -->')
8 plt.legend()
9 plt.show()
```



In [43]:

```

1 # Comparativo dos três modelos
2 print("EXTRA TREES CLASSIFIER:")
3 print("Accuracy Score:", accuracy_score(yteste, Train_predict_et))
4 print(classification_report(yteste, Train_predict_et))
5 print("REGRESSÃO LOGÍSTICA:")
6 print("Accuracy Score:", accuracy_score(yteste, Train_predict_lr))
7 print(classification_report(yteste, Train_predict_lr))
8 print("LIGHTGBM:")
9 print("Accuracy Score:", accuracy_score(yteste, Train_predict_lightgbm))
10 print(classification_report(yteste, Train_predict_lightgbm))

```

EXTRA TREES CLASSIFIER:

Accuracy Score: 0.7457716701902748

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| False | 0.70 | 0.86 | 0.77 | 1892 |
| True | 0.82 | 0.63 | 0.71 | 1892 |
| accuracy | | | 0.75 | 3784 |
| macro avg | 0.76 | 0.75 | 0.74 | 3784 |
| weighted avg | 0.76 | 0.75 | 0.74 | 3784 |

REGRESSÃO LOGÍSTICA:

Accuracy Score: 0.7618921775898521

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| False | 0.76 | 0.77 | 0.76 | 1892 |
| True | 0.76 | 0.76 | 0.76 | 1892 |
| accuracy | | | 0.76 | 3784 |
| macro avg | 0.76 | 0.76 | 0.76 | 3784 |
| weighted avg | 0.76 | 0.76 | 0.76 | 3784 |


LIGHTGBM:

Accuracy Score: 0.7418076109936576

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| False | 0.74 | 0.75 | 0.74 | 1892 |
| True | 0.75 | 0.73 | 0.74 | 1892 |
| accuracy | | | 0.74 | 3784 |
| macro avg | 0.74 | 0.74 | 0.74 | 3784 |
| weighted avg | 0.74 | 0.74 | 0.74 | 3784 |

In [44]:

```
1 # Comparando AUC dos modelos ExtraTrees, LogisticRegression e LightGBM
2 plt.figure(figsize=(5, 5), dpi=100)
3 plt.plot(logistic_fpr, logistic_tpr, marker='.', label='Logistic (auc = %0.3f)' % auc_)
4 plt.plot(et_fpr, et_tpr, marker='.', label='Extra Trees Classifier (auc = %0.3f)' % auc_)
5 plt.plot(lightgbm_fpr, lightgbm_tpr, marker='.', label='Light GBM (auc = %0.3f)' % auc_)
6 plt.xlabel('False Positive Rate -->')
7 plt.ylabel('True Positive Rate -->')
8 plt.legend()
9 plt.show()
```



A exigência da PUC Minas é de, no mínimo, três modelos de ML. Conforme demonstramos abaixo, o RandomForest gera basicamente o mesmo resultado do ExtraTrees. Dessa forma, utilizaremos, na versão a ser apresentada, os modelos ExtraTrees, LogisticRegression e LightGBM.

In [40]:

```

1 #Modelo RandomForest
2 rfm = RandomForestClassifier()
3 rfm = rfm.fit(xtreinamento, ytreinamento)
4 tp_rfm = rfm.predict(xteste)
5 print("Accuracy Score:", accuracy_score(yteste, tp_rfm))
6 print(classification_report(yteste, tp_rfm))

```

Accuracy Score: 0.7370507399577167


| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| False | 0.69 | 0.88 | 0.77 | 1892 |
| True | 0.83 | 0.60 | 0.69 | 1892 |
| accuracy | | | 0.74 | 3784 |
| macro avg | 0.76 | 0.74 | 0.73 | 3784 |
| weighted avg | 0.76 | 0.74 | 0.73 | 3784 |

In [41]:

```


1 skplt.metrics.plot_confusion_matrix(yteste, tp_rfm, normalize=True)
2 plt.show()

```



In [42]:

```
1 y_pred_rfm = rfm.predict_proba(xteste)
2 rfm_fpr,rfm_tpr,thereshold = roc_curve(yteste,y_pred_et[:,1])
3 auc_rfm = auc(rfm_fpr,rfm_tpr)
4 plt.figure(figsize=(5, 5), dpi=100)
5 plt.plot(rfm_fpr,rfm_tpr, marker='.', label='Extra Trees Classifier (auc = %0.3f)' % auc_rfm)
6 plt.xlabel('False Positive Rate -->')
7 plt.ylabel('True Positive Rate -->')
8 plt.legend()
9 plt.show()
```



Apêndice: testes mostrando variações que não foram aproveitadas no ML acima

1) Tentativa sem a coluna CEP

In []:

```
1 df_sem_cep = df_ml1[['Regiao', 'UF', 'CNAE', 'Atendida', 'CodAssunto', 'SexoConsumidor', 'FaixaEtaria', 'InscritoDAU']]
```

In []:

```
1 df_sem_cep
```

In []:

```
1 df_sem_cep = pd.get_dummies(df_sem_cep[['Regiao',
2                                     'UF',
3                                     'CNAE',
4                                     'Atendida',
5                                     'CodAssunto',
6                                     'SexoConsumidor',
7                                     'FaixaEtaria', 'InscritoDAU']])
```

In []:

```
1 df_sem_cep.shape # a do ml1 tinha 6922 colunas
```

In []:

```
1 Xdf_sem_cep = df_sem_cep.drop(['Atendida'],axis=1)
2 ydf_sem_cep = df_sem_cep.Atendida
3 smt = SMOTE()
4 Xdf_sem_cep_os,ydf_sem_cep_os = smt.fit_sample(Xdf_sem_cep,ydf_sem_cep) #os de oversampling
5 counter = Counter(ydf_sem_cep_os)
6 print(counter)
```

In []:

```
1 #Train_test_split nessa oversampled
2 # especificamos o tamanho do test_size = 0.3 pq assim as True/False do ytreinamento e ytestedf_sem_cep
3 xtreinamentodf_sem_cep, xtestedf_sem_cep, ytreinamentodf_sem_cep, ytestedf_sem_cep = train_test_split(Xdf_sem_cep_os, ydf_sem_cep_os, test_size=0.3, random_state=42)
```

ML do df_sem_cep

In []:

```
1 #Retomamos os train_test_split a partir do oversample ("_os") que já tínhamos feito
2 #xtreinamentodf_sem_cep, xtestedf_sem_cep, ytreinamentodf_sem_cep, ytestedf_sem_cep= train_test_split(Xdf_sem_cep_os, ydf_sem_cep_os, test_size=0.3, random_state=42)
```

In []:

```

1 # Modelo Extra Trees Classifier
2 etdf_sem_cep = ExtraTreesClassifier(random_state=0)
3 etdf_sem_cep = etdf_sem_cep.fit(xtreinamentodf_sem_cep, ytreinamentodf_sem_cep)
4 print("Acurácia de treinamento: ", etdf_sem_cep.score(xtreinamentodf_sem_cep, ytreinamentodf_sem_cep))
5 Train_predict_etdf_sem_cep = etdf_sem_cep.predict(xtestodf_sem_cep)
6 print("Acurácia de previsão: ", accuracy_score(ytestodf_sem_cep, Train_predict_etdf_sem_cep))
7 print(classification_report(ytestodf_sem_cep, Train_predict_etdf_sem_cep))

```

In []:

```

1 # Modelo Logistic Regression
2 lrdf_sem_cep = LogisticRegression(random_state=0)
3 lrdf_sem_cep = lrdf_sem_cep.fit(xtreinamentodf_sem_cep, ytreinamentodf_sem_cep)
4 print("Acurácia de treinamento: ", lrdf_sem_cep.score(xtreinamentodf_sem_cep, ytreinamentodf_sem_cep))
5 Train_predict_lrdf_sem_cep = lrdf_sem_cep.predict(xtestodf_sem_cep)
6 print("Acurácia de previsão: ", accuracy_score(ytestodf_sem_cep, Train_predict_lrdf_sem_cep))
7 print(classification_report(ytestodf_sem_cep, Train_predict_lrdf_sem_cep))

```

In []:

```

1 # Modelo Light Gradient Boosting Machine
2 import lightgbm
3 from lightgbm import LGBMClassifier
4 lightgbmdf_sem_cep = LGBMClassifier(random_state=0)
5 lightgbmdf_sem_cep = lightgbmdf_sem_cep.fit(xtreinamentodf_sem_cep, ytreinamentodf_sem_cep)
6 print("Acurácia de treinamento: ", lightgbmdf_sem_cep.score(xtreinamentodf_sem_cep, ytreinamentodf_sem_cep))
7 Train_predict_lightgbmdf_sem_cep = lightgbmdf_sem_cep.predict(xtestodf_sem_cep)
8 print("Acurácia de previsão: ", accuracy_score(ytestodf_sem_cep, Train_predict_lightgbmdf_sem_cep))
9 print(classification_report(ytestodf_sem_cep, Train_predict_lightgbmdf_sem_cep))

```

2) Tentativa com LabelEncoder para as de alta cardinalidade

A melhor técnica dispõe que devemos criar label encoders para as variáveis de alta cardinalidade. Não obstante, as nossas variáveis categóricas não podem sofrer o enviesamento: um CNAE 8630502 não é 8630502 vezes melhor ou pior do que o CNAE 0000001. Adotar os label encoders para as variáveis verdadeiramente explicativas - CNAE e CodAssunto - implicaria em um descolamento com a realidade dos problemas dessa natureza.

In []:

```
1 df_le = df_m11 #leia-se "dataframe label encoder"
```

In []:

```
1 from sklearn.preprocessing import LabelEncoder
```

In []:

```
1 df_le.nunique()
```

In []:

```
1 colunascategoricas = df_le.select_dtypes('object').columns
2 colunascategoricas
```

In []:

```
1 for col in colunascategoricas:
2     df_le[col+'_encoded'] = LabelEncoder().fit_transform(df_le[col])
```

In []:

```
1 df_le
```

In []:

```
1 # Para definir as colunas que permanecerem no nosso dataframe selecionaremos:
2 # para de alta cardinalidade, mantemos LabelEncoder;
3 # para as de baixa e média, criaremos dummies
```

In []:

```
1 df_le = df_le[['Regiao', 'UF', 'CNAE_encoded', 'Atendida', 'CodAssunto_encoded', 'SexoConsumidor', 'FaixaEtaria']]
```

In []:

```
1 df_le
```

In []:

```
1 # Pois bem, vamos agora criar dummies para 'Regiao' (5 unique), 'UF' (15), SexoConsumidor (2), FaixaEtaria (5)
```

In []:

```
1 df_le_du = pd.get_dummies(df_le,
2                             columns = ['Regiao',
3                                         'UF',
4                                         'SexoConsumidor',
5                                         'FaixaEtaria'],
6                             prefix = ['Regiao',
7                                       'UF',
8                                       'SexoConsumidor',
9                                       'FaixaEtaria'],
10                            prefix_sep = '_' )
11
```

In []:

```
1 df_le_du.info()
```

In []:

```
1 df_le_du # leia-se df label encoder com dummies
```

In []:

```
1 df_le_du.info()
```

Agora, então, temos dummies para as de baixa e média cardinalidade: 'Regiao' (5 unique), 'UF' (15 unique), 'SexoConsumidor' (2 unique) e 'FaixaEtaria' (7); e label encoder para as de alta cardinalidade 'CNAE_encoded' (368 unique), 'CodAssunto_encoded'(175 unique) e CEP (6354 unique)

In []:

```
1 Xdf_le_du = df_le_du.drop(['Atendida'],axis=1)
2 ydf_le_du = df_le_du.Atendida
3 smt = SMOTE()
4 Xdf_le_du_os,ydf_le_du_os = smt.fit_sample(Xdf_le_du,ydf_le_du) #os de oversampled
5 counter = Counter(ydf_le_du_os)
6 print(counter)
```

In []:

```
1 #Train_test_split nessa oversampled
2 # especificamos o tamanho do test_size = 0.3 pq assim as True/False do ytreinamento ficam
3 #True/False do yteste também ficam iguais
4 xtreinamentodf_le_du, xtestedf_le_du, ytreinamentodf_le_du, ytestedf_le_du = train_test_split(Xdf_le_du_os, ydf_le_du_os, test_size=0.3, random_state=42)
```

In []:

```
1 #Retomamos os train_test_split a partir do oversample ("_os") que já tínhamos feito
2 #xtreinamentodf_le_du, xtestedf_le_du, ytreinamentodf_le_du, ytestedf_le_du = train_test_split(Xdf_le_du_os, ydf_le_du_os, test_size=0.3, random_state=42)
```

ML do df_le_du ("df label encoder dummies")

In []:

```
1 # Modelo Extra Trees Classifier
2 etdf_le_du = ExtraTreesClassifier(random_state=0)
3 etdf_le_du = etdf_le_du.fit(xtreinamentodf_le_du, ytreinamentodf_le_du)
4 print("Acurácia de treinamento: ", etdf_le_du.score(xtreinamentodf_le_du, ytreinamentodf_le_du))
5 Train_predict_etdf_le_du = etdf_le_du.predict(xtestedf_le_du)
6 print("Acurácia de previsão: ", accuracy_score(ytestedf_le_du, Train_predict_etdf_le_du))
7 print(classification_report(ytestedf_le_du, Train_predict_etdf_le_du))
```

In []:

```
1 # Modelo Logistic Regression
2 lrdf_le_du = LogisticRegression(random_state=0)
3 lrdf_le_du = lrdf_le_du.fit(xtreinamentodf_le_du, ytreinamentodf_le_du)
4 print("Acurácia de treinamento: ", lrdf_le_du.score(xtreinamentodf_le_du, ytreinamentodf_le_du))
5 Train_predict_lrdf_le_du = lrdf_le_du.predict(xtestedf_le_du)
6 print("Acurácia de previsão: ", accuracy_score(ytestedf_le_du, Train_predict_lrdf_le_du))
7 print(classification_report(ytestedf_le_du, Train_predict_lrdf_le_du))
```

In []:

```
1 #Modelo Light Gradient Boosting Machine
2 import lightgbm
3 from lightgbm import LGBMClassifier
4 lightgbmdf_le_du = LGBMClassifier(random_state=0)
5 lightgbmdf_le_du = lightgbmdf_le_du.fit(xtreinamentodf_le_du, ytreinamentodf_le_du)
6 print("Acurácia de treinamento: ", lightgbmdf_le_du.score(xtreinamentodf_le_du, ytreinamentodf_le_du))
7 Train_predict_lightgbmdf_le_du = lightgbmdf_le_du.predict(xtestodf_le_du)
8 print("Acurácia de previsão: ", accuracy_score(ytestodf_le_du, Train_predict_lightgbmdf_le_du))
9 print(classification_report(ytestodf_le_du, Train_predict_lightgbmdf_le_du))
```

Não obstante melhorar o desempenho dos modelos Extra Trees e LightGBM, houve piora no modelo escolhido, o de Regressão Logística.